

SPOONS: NETFLIX OUTAGE DETECTION USING MICROTTEXT
CLASSIFICATION

A Thesis

Presented to

the Faculty of California Polytechnic State University

San Luis Obispo

In Partial Fulfillment

of the Requirements for the Degree

Master of Science in Computer Science

by

Eriq Augustine

March 2012

© 2012

Eriq Augustine

ALL RIGHTS RESERVED

COMMITTEE MEMBERSHIP

TITLE: SPOONS: Netflix Outage Detection Using
Microtext Classification

AUTHOR: Eriq Augustine

DATE SUBMITTED: March 2012

COMMITTEE CHAIR: Alex Dekhtyar, Ph.D.

COMMITTEE MEMBER: Clint Staley, Ph.D.

COMMITTEE MEMBER: Franz Kurfess, Ph.D.

COMMITTEE MEMBER: Foaad Khosmood, Ph.D.

Abstract

SPOONS: Netflix Outage Detection Using Microtext Classification

Eriq Augustine

Every week there are over a billion new posts to Twitter services and many of those messages contain feedback to companies about their services. One company that has recognizes this unused source of information is Netflix. That is why Netflix initiated the development of a system that lets them respond to the millions of Twitter and Netflix users that are acting as sensors and reporting all types of user visible outages. This system enhances the feedback loop between Netflix and its customers by increasing the amount of customer feedback that Netflix receives and reducing the time it takes for Netflix to receive the reports and respond to them.

The goal of the SPOONS (Swift Perceptions of Online Negative Situations) system is to use Twitter posts to determine when Netflix users are reporting a problem with any of the Netflix services. This work covers the architecture SPOONS system and framework as well as outage detection using tweet classification.

Acknowledgements

Thanks Alex, ABRA, Netflix especially Kevin McEntee, and all the funions.
Thanks to Farscape for encouraging team bonding and providing a common enemy.

Contents

List of Tables	ix
List of Figures	x
1 Introduction	1
1.1 General Problem: Swift Perception Of Online Negative Situations	1
1.2 Solution Overview	3
1.3 Ethics of Twitter Observation	5
1.3.1 Twitter Terms of Service	5
1.4 SPOONS Requirements	6
1.5 Contributions and Organization	8
2 Background & Related Work	9
2.1 Text Stream Analysis	9
2.2 Classification	9
2.2.1 Microtext Classification	9
2.2.2 WEKA	9
2.3 Twitter API	10
2.3.1 Rate Limiting	10
2.3.2 Pagination	10
2.3.3 Query Anatomy	11
2.3.4 Result Anatomy	12
3 SPOONS Architecture	14
3.1 Framework Architecture	14
3.1.1 High Level Solution	14

3.1.2	Gatherers	17
3.1.3	Processors	18
3.1.4	Analysis Pipelines	19
3.1.5	Tasks	20
3.1.6	Modelers	20
3.1.7	Monitors	21
3.1.8	Control	21
3.1.9	Distributed Model	23
3.2	Database	27
3.2.1	Tables and Schemas	28
3.2.2	UI Stored Procedures	30
4	Classifiers	31
4.1	Fitting Into The SPOONS Framework	31
4.2	WEKA Classifiers	31
4.3	Non-WEKA Classifiers	32
4.3.1	BPNB	32
4.4	Feature Selection	32
4.4.1	Text Filtering	32
4.5	Tweet Classes	36
4.5.1	Tweet Groups	36
4.6	Training Set	37
5	Experimental Setup	39
5.1	Ground Truth	39
5.2	Success Metrics	39
5.3	Classifier Evaluation	40
5.4	Outage Detection	40
5.4.1	Monitor Parameters	41
6	Results	42
6.1	Classifier Evaluation	42
6.2	Outage Detection	42

7	Conclusions	59
7.1	Current Limitations of SPOONS	59
7.2	Current and Future Work	59
7.2.1	Advanced Sentiment Analysis	59
	Glossary	59
	Bibliography	60

List of Tables

4.1	Netflix-related Twitter Traffic	38
6.1	Uncompressed, NGram, NoFilter Classification Results	43
6.1	Uncompressed, NGram, NoFilter Classification Results Cont. . . .	44
6.1	Uncompressed, NGram, NoFilter Classification Results Cont. . . .	45
6.2	Uncompressed, NGram, EriqFilter Classification Results	46
6.2	Uncompressed, NGram, EriqFilter Classification Results Cont. . .	47
6.2	Uncompressed, NGram, EriqFilter Classification Results Cont. . .	48
6.3	Uncompressed, TweetFSG, NoFilter Classification Results	49
6.3	Uncompressed, TweetFSG, NoFilter Classification Results Cont. .	50
6.3	Uncompressed, TweetFSG, NoFilter Classification Results Cont. .	51
6.4	Uncompressed, TweetFSG, EriqFilter Classification Results	52
6.4	Uncompressed, TweetFSG, EriqFilter Classification Results Cont.	53
6.4	Uncompressed, TweetFSG, EriqFilter Classification Results Cont.	54
6.5	Compressed, NGram, NoFilter Classification Results	55
6.6	Compressed, NGram, EriqFilter Classification Results	56
6.7	Compressed, TweetFSG, NoFilter Classification Results	57
6.8	Compressed, TweetFSG, EriqFilter Classification Results	58

List of Figures

1.1	Outage Tweets Example	4
1.2	System Concept Diagram	4
2.1	WEKA logo	10
2.2	Twitter Search API Result	13
3.1	SPOONS Framework Architecture	15
3.2	SPOONS UI	15
3.3	Twitter Holes	18
3.4	SPOONS Server Architecture	23
3.5	SPOONS Distributable Task Flow	24
4.1	SPOONS Groups	37

Chapter 1

Introduction

1.1 General Problem: Swift Perception Of On-line Negative Situations

Twitter is an immensely popular micro-blogging service. According to Twitter as of March 14th 2011, approximately one billion micro-posts, *tweets*, were being posted per week[18]. Because of the low time and effort cost of tweeting, only a few seconds from a smart phone, users of Twitter post tweets about almost every aspect of their daily lives. Because of this large stream of information, Twitter makes an excellent source of information for data miners. Already, researchers have been using Twitter to attempt to track and model disease outbreaks[5], earthquakes[12], and the stock market[11].

Netflix is the one of the largest online Internet subscription service for streaming movies and television shows. Netflix has over 25 million subscribers watching media streamed to over 450 different platforms. Even a short disruption of service can effect millions of users. Therefore, quickly detecting service outages is

essential to keep customers happy. However, service outage detection is no trivial matter in Netflix's environment. In addition to constantly streaming thousands of different videos to hundreds of different platforms, Netflix also has to deal with problems caused by their entire infrastructure being hosted in the cloud with Amazon's AWS.

Netflix saw the power in Twitter as a potential data source for detecting service outages that is orthogonal to their current, more traditional outage detection methods. Currently, Netflix utilizes four different methods for detecting outages:

Internal Monitoring Systems. Like any sizable service providing company, Netflix utilizes many different internal monitoring systems to detect service outages. However, there are some class of problems that are difficult to solve with internal monitoring. These problems include corrupt video files or a problem on a third-party delivery platform such as Roku or AppleTV. These problems are obvious to the end user, but very difficult to detect internally. In addition, the internal monitoring systems share the same infrastructure as the service providing system. Therefore, a problem in the infrastructure can cause both systems to go down at the same time.

External Monitoring Systems. Netflix contracts with external services that can periodically probe its systems to try and detect problems. However, this model too has problems. There are many problems that cannot be seen from an external probe. Also, if this system probes too often then it is taking compute time away from the servers that are trying to deliver content to end users.

Customer Service. Calls to customer service are a very straight-forward way to detect outages. Unfortunately, this method is very slow and inconsistent. It takes a lot of frustration to get a user to lookup a phone number and complain.

Manual Twitter Observation. Manual observation shows that there is usually a response on Twitter when Netflix suffers a service outage. Figure 1.1 shows some tweets that occurred during a disruption of Netflix’s service to the Nintendo Wii. However without any infrastructure, Twitter observation is slow and inconsistent. It is also very time consuming to have someone constantly watching Twitter for signs of an outage.

Given all these deficiencies Netflix wanted a monitoring system that is separate from their infrastructure, fast, and does not require any human intervention[13].

1.2 Solution Overview

SPOONS (Swift Perception Of Online Negative Situations) is a system that is designed to use tweets to detect outages in Netflix content delivery systems. At present the system supports a wide variety of detection methods that use some combination of time series analysis, classification, natural language processing, sentiment analysis, and filtering.

Figure 1.2 shows how the SPOONS system can be divided into three main parts: **input**; **analysis methods**; and **output**. The inputs are tweets gathered from Twitter. Then the analysis methods use a combination of sentiment estimation, classification, and traffic volume analysis to detect when an outage is occurring. The outputs of the system are: email alerts to Netflix engineers, and a web UI that displays information about the outage.

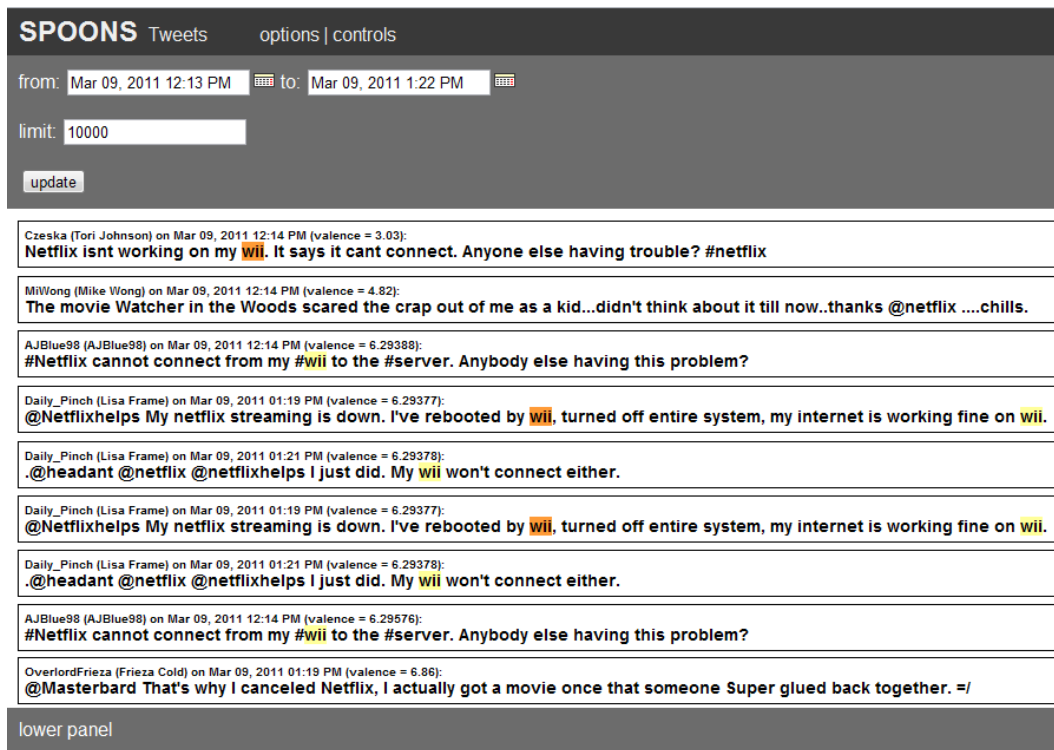


Figure 1.1: Tweets posted on March 9, 2011 during a disruption of Netflix streaming to the Nintendo Wii console.




Figure 1.2: This system concept diagram shows the general flow of processing done in the SPOONS system.

1.3 Ethics of Twitter Observation

The work in this project uses content that users post on Twitter without their knowledge. This monitoring system isn't being announced to the public because widespread knowledge of it would increase the likelihood of a malicious attack. This practice may lead to concerns about the level of privacy or ownership being provided to Twitter users regarding the content they post through the Twitter services. The goal of this section is to address these concerns by providing more information about the Twitter services and how the SPOONS system and this work uses the tweets.

1.3.1 Twitter Terms of Service

According to Twitter Terms of Service[19] agreement that everyone accepts automatically by accessing or using Twitter services:

“You retain your rights to any Content you submit, post or display on or through the Services. By submitting, posting or displaying Content on or through the Services, you grant us a worldwide, non-exclusive, royalty-free license (with the right to sublicense) to use, copy, reproduce, process, adapt, modify, publish, transmit, display and distribute such Content in any and all media or distribution methods (now known or later developed).”

“This license is you authorizing us to make your Tweets available to the rest of the world and to let others do the same.”

“You agree that this license includes the right for Twitter to make such Content available to other companies, organizations or individuals who partner with Twitter for the syndication, broadcast, distribution or publication of such Con-

tent on other media and services, subject to our terms and conditions for such Content use.”

“We encourage and permit broad reuse of Content. The Twitter API exists to enable this.”

“Such additional uses by Twitter, or other companies, organizations or individuals who partner with Twitter, may be made with no compensation paid to you with respect to the Content that you submit, post, transmit or otherwise make available through the Services.”

In short, Twitter takes ownerships of users’ tweet as soon as they are posted on Twitter. Using the Twitter API allows SPOONS to obtain the tweets with the consent of Twitter. Therefore, the collection and analysis of Twitter data by SPOONS is well withing the Twitter Terms of Service.

1.4 SPOONS Requirements

Netflix has provided the following set of key requirements to be met by the SPOONS system:

Structural Independence. The outage detection system shall be structurally independent of both the software and the hardware infrastructure used by Netflix. It shall rely only on information that is publicly available and free for use. This ensures that the outage detection system stays up even when any or all Netflix servers are experiencing downtime.

Use of Amazon Web Services. Netflix is one of the largest customers of Amazon.com’s cloud computing service, Amazon Web Services (AWS). AWS allows users to create new cloud machines (instances) in many regions throughout the world. The outage detection system shall be deployed on one or more AWS servers that are operationally independent of other AWS servers used by Netflix. Using a cloud solution allows the outage detection and alert system to be deployable on a global scale.

Real-Time. Netflix’s streaming services run in real-time and any downtime has an immediate impact on customers. To minimize that impact, the outage detection system shall notify Netflix of detected outages as soon as possible.

Precise Outage Detection. The number of non-outage situations that raise an alert shall be minimized. While a small number of false positives detected in real-time may be acceptable, the outage detection system shall detect outages and generate alerts with as high precision as possible.

Comprehensive Outage Detection. Not all Netflix outages will generate a signal on Twitter. Those that don’t may be allowed to go unnoticed by the outage detection system (as the system will have no basis for detecting them), but any outage that causes a signal on Twitter shall be detected.

User-Friendly Online UI. The outage detection and alert system shall have an easy-to-use, informative, online UI which shall provide Netflix employees with real-time information and historic data about the state of Netflix according to Twitter. The information provided shall include:

- times of outages;
- times of other anomalous events;
- current and recent Netflix-related Twitter traffic trends;
- and samples of Netflix-related tweets.

1.5 Contributions and Organization

SPOONS is a continual team effort and has been touched and improved by many different people. The idea originated at Netflix and was passed to the ABRA team at Cal Poly. The ABRA team has published a paper on SPOONS [2]. In addition, Cailin Cushing defended a thesis using SPOONS[6].

The main contributions of this work are the design and implementation of the SPOONS system, framework, server architecture, distribution model, and database schema. As well as the design and implementation of classification based outage detection methods.

The rest of the paper is organized as follows. Chapter 2 covers background and related work. Chapter 3 discusses the architecture of SPOONS. Chapter 4 discuss the details of the classifiers used in SPOONS. Chapters 5 and 6 respectively discuss the different experiments performed and their results. Chapter 7 wraps up the paper.

Chapter 2

Background & Related Work

2.1 Text Stream Analysis

Text Stream Analysis [3][7][10].

2.2 Classification

Classification [14]

2.2.1 Microtext Classification

Microtext Classification [9]

2.2.2 WEKA

SPOONS utilizes several classifiers provided in the *WEKA Machine Learning Package*. WEKA is an open source package written under the GNU General



Figure 2.1: WEKA logo

Public License[8].

2.3 Twitter API

All of the data data that SPOONS uses is obtained in real time using the Twitter Search REST API[20].

2.3.1 Rate Limiting

Twitter imposes a limit on the number of queries to the Search API. However, Twitter does not publish the official limit. However, our experiments suggest that SPOONS can query the API for all new Tweets once every two minutes without suffering from rate limiting.

2.3.2 Pagination

Twitter paginates the results from its search API. The maximum results you can get per page is 100, and each query can return at most 15 pages. Therefore when there are more than 1500 tweets generated per minute, SPOONS must do multiple search queries.

2.3.3 Query Anatomy

Our typical Twitter API query is structured as follows:

```
http://search.twitter.com/search.json?  
q=<query>&rpp=100&result_type=recent&  
since_id=<tweet id>&max_id=<tweet id>
```

The parameters are:

json: Twitter can supply the result data in either ATOM or JSON format. Testing with both have shown that the ATOM results are less consistent and provide less data. Because of the more accurate information returned from the JSON API, we are able to write more efficient queries. Using the ATOM API, we could query Twitter only once every five minutes; as opposed to every two minutes with the JSON API.

q: The search query. Twitter supports some advanced search features such as conjunction and negation.

rpp: “Results Per Page”. Twitter paginates the responses from the Search API. SPOONS always uses the maximum pagination value to decrease the number of requests per hour and lessen the chance of being rate limited.

result_type: Twitter allows users to get results ordered by either relevance or time. Since we want to gather all tweets about our query, we choose to get the results ordered by time. In addition, the “since_id” and “max_id” parameters do not work when results are sorted by relevance.

since_id: The id of the oldest tweet that should be returned. This is not a hard limit, but provides a nice starting point.

max_id: The id of the most recent tweet that should be returned. It may seem counter-intuitive to provide a cap on the most recent tweet, when one wants to query for all of the most recent tweets. However when a query spans across more than 15 pages, it will need to be broken into a new query restarting at the first page. In this situation, not providing an upper limit will include new tweets outside of the original search scope. This can result in tweets are forever lost to us.

2.3.4 Result Anatomy

Figure 2.2 shows the result from the query “eriq netflix”. Notice that some fields, like the `geo` field, can be null. Also note that the API incorrectly guessed the language of the tweet as Danish.

```

{
  completed_in: 0.012,
  max_id: 298199940868489200,
  max_id_str: "298199940868489216",
  page: 1,
  query: "netflix+eriq",
  refresh_url: "?since_id=298199940868489216&q=netflix%20eriq&result_type=recent",
  results: [
    - {
      created_at: "Sun, 03 Feb 2013 22:43:00 +0000",
      from_user: "eriq_augustine",
      from_user_id: 238374031,
      from_user_id_str: "238374031",
      from_user_name: "eriq",
      geo: null,
      id: 298199940868489200,
      id_str: "298199940868489216",
      iso_language_code: "da",
      metadata: {
        result_type: "recent"
      },
      profile_image_url: "http://a0.twimg.com/sticky/default_profile_images/default_profile_0_normal.png",
      profile_image_url_https: "https://si0.twimg.com/sticky/default_profile_images/default_profile_0_normal.png",
      source: "&lt;a href=&quot;http://twitter.com/&quot;&gt;web&lt;/a&gt;",
      text: "eriq love netflix",
      to_user: null,
      to_user_id: 0,
      to_user_id_str: "0",
      to_user_name: null
    },
  ],
  results_per_page: 100,
  since_id: 0,
  since_id_str: "0"
}

```

Figure 2.2: A JSON result from the Twitter Search API

Chapter 3

SPOONS Architecture

There are multiple levels of architecture within SPOONS that need to be discussed. There is the Framework Architecture (Figure 3.1) that describes the relations between the different pieces of the framework; the Server Architecture (Figure 3.4) that describes the layout of the different servers involved in the SPOONS system; and the Distribution Model which describes how tasks are distributed between the different servers.

3.1 Framework Architecture

This section describes the architecture of the SPOONS framework. The SPOONS framework includes all pieces of SPOONS that take the data from gathering all the way through to final analysis.

3.1.1 High Level Solution

The general solution taken by SPOONS consists of four main steps:




Figure 3.1: The flow of control and data through the SPOONS framework system.




Figure 3.2: The web UI for SPOONS.

Collect Tweets are collected from Twitter.

Process The tweets are converted from plain text to some form of information that can be analyzed.

Model Use the information generated from the previous step to build a mathematical model of the information. Use past information to predict what the current model of the data should look like.

Compare Compare the two models generated in the previous step. A significant divergence means that there is anomalous traffic.

Framework Overview

Figure 3.1 shows the flow of control and data through the SPOONS framework. Data comes into SPOONS in the form of Tweets collected by the Gatherers, and leave SPOONS in the form of alerts generated by the Monitors.

Gatherer. Gatherers are responsible for collecting documents from a specified data source such as the Twitter Search API.

Control. The Control is responsible for controlling the SPOONS server. It maintains data structures with all of the Gatherers and Analysis Pipelines. It is also responsible for communication with other servers in the SPOONS cluster.

Processor. Processors are data transformation utilities that takes raw data and puts it in a form that other components can use.

Modeler. Modelers are responsible for building a mathematical model of the data and can be split into two groups: **Predictors** and **Counters**. Predictors build a predictive model of the data. Counters to build a model of the data that was actually gathered by the system.

Monitor. Monitors take the models produced by the Predictors and Counters and compares them. The Monitors are responsible for making the final decision on about a period of time being anomalous.

3.1.2 Gatherers

The data enters SPOONS at the Gatherers. The Gatherers run periodically (for Twitter, every two minutes). Gatherers are asynchronous and not dependent on any other part of the framework. There may be multiple different Gatherers running on the same machine. Gatherers are abstracted to be able to gather data from any source. Once the Gatherers get their data, they place the data in the database and notify the Control that there is new data available to the system.

Twitter Holes

It is worth noting that sometimes the Twitter Search API fails to return any data. We have not discovered the cause of this, but Twitter does not report any errors. For unspecified amounts of time the Twitter API will just report zero new tweets. We call these dead zones “holes”. We have found that a query from a different IP usually does not experience the same hole. To counteract holes, we run Gatherers on multiple servers and resolve uniques upon insertion into the database.




Figure 3.3: One server in a hole is covered by two other gathering servers.

3.1.3 Processors

Processors are responsible for processing or transforming data before it goes into the analysis pipelines.

- **Classifier Processors** There exists a Processor for every tweet classifier used in SPOONS (see Chapter 4. Because of the high number of classifiers used, these constitute the majority of Processors and the largest unit of work in SPOONS. These Processors classify every tweet into one of the nine tweet categories discussed in Section 4.5.
- **Author Processors** The Author Processors extract the author of tweets and try to establish which authors are credible. These Processors are outside the scope of this work and are discussed in other work[?].
- **Valence Processors** The Valence Processors assign a numeric “happiness” score to every tweet. How that score is produced is outside the scope of this work (see Section 7.2.1).

- **Document Frequency Processors** The Document Frequency Processors maintain term frequencies and inverse document frequencies for the collection of tweets in SPOONS.

Unlike most parts of the analysis pipeline, Processors are a shared resource. That is, multiple analysis pipelines invoke the same Processors. However, it does not make sense to restart the processing once it is started, or to start another instance of the same Processor for the same data. Processors have a finite amount of data to process and may be cumulative. To make sure that no redundant work is done, Processors are singleton. When multiple threads call into a Processor to do work, the Processor will block all incoming threads until the work is complete. Then, the Processor will release all of the threads that requested the work. This model allows all the analysis pipelines to share the same Processor without any redundancies.

3.1.4 Analysis Pipelines

An Analysis Pipeline (also called Analysis Method) is the analytical center of the SPOONS framework. The pipeline aggregates multiple tasks that it needs to run on the data.

An Analysis Pipeline typically starts with running any number of Processors on the data. Then, the pipeline invokes modelers on the data from the Processors. These modelers typically build models for the actual data coming into SPOONS as well as predictive models. Finally, the pipeline invokes tasks that assess the models produced in the previous step and decides whether or not there is an anomaly.

Every Analysis Pipeline gets its own thread, and there is no interdependence

between the different pipelines. Currently, SPOONS usually runs more than 20 Analysis Pipelines at a time.

3.1.5 Tasks

Tasks are the core unit of computation in SPOONS. Almost everything that can be “run” is a child of the Task base class. Every Task gets its own thread, and callers into the Task may request that the task block the calling thread until the Task is complete.

Tasks are singleton with respects to the leaf child class. Therefore there are many tasks, but every task is unique. We do this by enforcing that the class name is unique upon construction. The uniqueness of tasks is very important to SPOONS distribution model that will be discussed in Section 3.1.9.

3.1.6 Modelers

Modelers are Tasks that are responsible for building a mathematical representation for the data.

Predictors

Predictors build a predictive model of the data. For example, we have noticed that tweet volume tends to be periodic day-to-day and week-to-week. Therefore, a Predictor may model that prediction by guessing that the volume in the future will be the same as it was the previous week or day.

Counters

Counters attempt to build a model of data that was actually gathered by the system. Going with the previous example, the Counter for modeling tweet volume would simply count the number of tweets gathered for a period.

3.1.7 Monitors

Monitors take the models produced by the Predictors and Counters and compares them. The Monitors are responsible for making the final decision about a period of time being anomalous.

Auto-Tuning

Monitors take anywhere from two to six tuning parameters. To find the best set of parameters, the Monitors can automatically run themselves on a training set and search the space of all possible parameters. They then keep the parameters that result in the best score.

3.1.8 Control

The Control is the center of a SPOONS instance. It handles the flow of all control and has the ability to start and stop any task or pipeline on demand. It holds references to all the threads for the Gatherers and Analysis Pipelines. The Control handles all the setup and tear down in the system.

There are different types of Controls that decide the behavior SPOONS on each respective server. The Control is singleton with respects to the base class.

Therefore, only one instance of any type of Control can be active at any given time.

The Control is very careful to never allow anyone to own a reference to the currently running Control. All requests to the Control are made statically to the “Control” base class. The base class will then forward the request onto the specific instance of Control. We do this so that the rest of the SPOONS system will never know what kind of Control is currently active. So we can switch a server between different roles without restarting the system or notifying any other components of the SPOONS system.

All Controls will always run the entire slew of Gatherers.

Master Control

The Master Control is the Control that is responsible for the controlling SPOONS when it is in distributed mode. The Master Control maintains information on all the active worker servers. It will send the worker servers messages to tell them what work to do.

The Master Control maintains “shallow execution” of every pipeline in the system. This means that this control will run each pipeline, but then distribute work for each pipeline as it is created.

Worker Control

Worker Controls do not take any initiative to run any tasks. Instead, they just wait for a Master Control to tell them what to do.




Figure 3.4: The server architecture of the SPOONS system.

Single Control

The Single Control is for a SPOONS instance that wants to run on a single server.

3.1.9 Distributed Model

As discussed before, SPOONS is a multi-server system (Fig 3.4). The SPOONS system uses the master/worker paradigm with a single master and N workers.

All of the servers share two primary resources: the primary database and a NoSQL property store. When a master or worker comes online, it inserts an entry for itself into the shared property store. If the new server is a worker, it will alert the current master about its existence; and visa-versa if the new server is a master. In addition, all workers are required to heartbeat to the master every

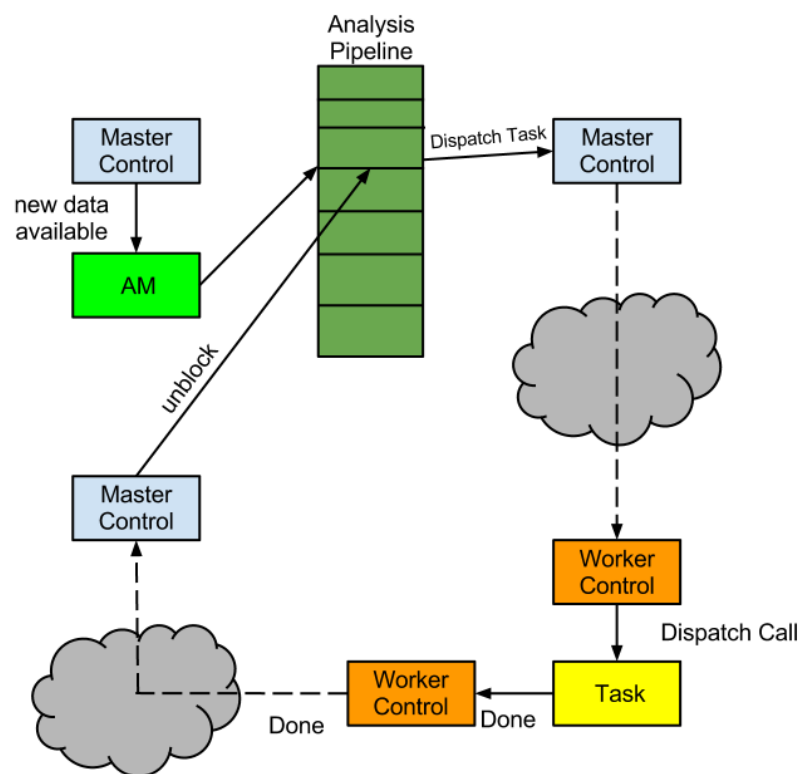


Figure 3.5: The control flow for distributable tasks.

15 seconds and the master heartbeats to the workers every 15 seconds. Using this system, the master always knows about all of the workers and the worker always knows about the current master. When a server misses three heartbeats, the server expecting that heartbeat assumes that the server has gone down.

Distribution Assumptions

The SPOONS distribution model relies on two assumptions about the system: every server contains exactly the same data in memory and every Task can be uniquely referenced.

Same Data. SPOONS assumes that every server will have the same data in memory on every server. This means that not only does every server need to have the same data structures in memory, but also that every server needs to have the same classes instantiated. The only exception to this assumption is the Control. Depending on the role of the server, a different Control will be instantiated. Because of this assumption, we do not have to worry about active replication between servers or a worker being asked to do work that requires a class that is not instantiated.

Uniquely Referenced Tasks. As stated in Sec 3.1.5, Tasks are the basic unit of work inside SPOONS. When a worker is told to execute some work, it is being asked to execute a specific task with specified parameters. Therefore, Tasks need to be able to be referenced by a key that can be serialized and sent over the wire from the master to the worker.

Distributable Tasks

Distributable Task is a subclass of Task that provides some of the distribution mechanism for Tasks. When a task is to be distributed, the Distributable Task calls into the Control and requests that the Control distributes it. The next step varies depending on the type of Control that is active:

Master Control. The task distributing control flow is described in Image ??¹. A Master Control will take the pause the calling thread and send a message to a selected worker¹ telling to run the task with given parameters. The message that goes to the worker just contains the task's unique identifier and the parameters to the task's run. When the task is complete, the Worker Control will send the Task's return status back to the Master Control. When the master receives a message from the worker that the requested task has completed its run, it will resume the original calling thread and have it return with the return status given by the worker.

Worker Control. If a Worker Control receives a task, then it is being asked to distribute a task that is already being distributed. We consider this a violation of the framework and will throw an error.

Single Control. A worker will just call back into the task and tell it to run itself.

¹The current scheduling algorithm chooses the worker that has the fewest tasks currently assigned to it.

Shared Properties

As previously stated, all servers must maintain a consistent in-memory view of the system. This can be troublesome if a Task needs to maintain cumulative settings or member datum. Not only will this data need to be consistent on all the servers, but it also needs to maintain this data between starts and stop of the system. An Analysis Pipeline should be able to the stopped for an arbitrary amount of time and then restarted without losing data or its place.

To enforce these restrictions, we use a shared property store. The shared property store is a MongoDB server. Whenever a Task needs to store member datum, it places it in the shared store. Therefore, any server may access this data. A Task can first be run on Server A and then on Server B. Because it stores the necessary information in the shared property store, Server B can have all the information gained from the run on Server A and not lose any positional information.

In addition to storing shared properties, the shared property store houses information on every active server. When a server comes online, it queries the property store to find all the other active servers and inserts itself into the store. If a server fails to heartbeat, then the servers that still live will remove the entry that server from the property store.

3.2 Database

SPOONS is backed by a MySQL database. SPOONS currently uses 225 tables and 35 stored procedures.

3.2.1 Tables and Schemas

Each stage in an analysis pipeline generally stores some information in the database. Because each stage generally deals with similar types of data, these tables are considered to be in the same group. We enforce group membership using hints in the table names. For example, the table name “**RESULT**_EN_class_heuristic_bayes_net” gives five hints as to the type of the table.

1. **RESULT** - Marks this table as a result table. This means that it is guaranteed to be shown in the UI.
2. EN - The language of the tweets that were input into this method.
3. class - Indicates that these results are output from a tweet classifier.
4. heuristic - States that the type of classifier used was a heuristic classifier.
5. bayes_net - The name of the classifier used.

Using all of these hints, the UI can then ask for data for specific types of tables (eg. all English table).

The six different top level categories that SPOONS recognizes are:

1. CALC - These are intermediary tables in analysis pipelines. They are never shown to the UI.
2. CONFIG - Contains information that analysis methods use to configure themselves before runs. These tables have been mostly replaced with the shared property store (see Section 3.1.9).
3. DATA - Raw input data. These tables are generally the output from the Gatherers.

4. META - Contains information that is not analyzed, but required to the system. For example, the different classes that the classifiers use along with descriptions of each class.
5. RESULT - These tables are output from some analysis pipeline. They are guaranteed to be shown in the UI.
6. TEST - These tables are used for debuggin and development. They will never be shown in a user-facing UI, however may be shown in development UIs.

Tweets Table

As the most used and important table in the database, the table that houses all of our tweets, “DATA_tweets”, gets special attention.

The tweets table contains ten rows:

- id - An auto-incremented primary key.
- twitter_id - The unique id Twitter gives a tweet.
- published - The epoch time that the tweet was posted according to Twitter.
- content - The raw content of the tweet.
- source - Information on where the tweet was posted from (eg. from a third party app).
- lang - The suggested language of the tweet.
- author - The author of the tweet.

- `frame_id` - The frame that this tweet falls into, has an index on it.
- `place` - Information on where the tweet was posted from. This is a JSON structure and may contain fields such as “city” and “state”.
- `geo` - Geographical coordinates of place.

Frames. Inside SPOONS, we use a “frame” as the atomic unit of time. Currently, a frame corresponds to a minute. Bucketing the tweets into frames allows us to gain a natural aggregation and smoothing. It also provides a natural index. Maintaining an index on *frame_id* allows quick retrieval of timeseries data which is the primary task of SPOONS. Because insertions are generally chronological, insertions are also quick and do not require a rebuild of the B-Tree index[1].

3.2.2 UI Stored Procedures

In addition to utility procedures, the database holds many stored procedures used by the UI. This keeps the UI fairly stable in the face of database changes.

Chapter 4

Classifiers

4.1 Fitting Into The SPOONS Framework

Although the classifiers can be used at any stage in an analysis pipeline, the classifiers are usually implemented as a Processor. It takes in a range of tweets and produces a classification for each tweet.

4.2 WEKA Classifiers

SPOONS uses several classifier from the WEKA machine learning package[8].

- Naive Bayes
- Bayes Net
- J48 (a method of generating a C4.5 decision tree[17])
- K-Nearest Neighbors

- SMO (Support Vector Machine trained with Sequential Minimal Optimization[15])

4.3 Non-WEKA Classifiers

In addition, I have implemented two classifiers by hand. WEKA contains many different types of classifiers, but it can be inflexible and can be inefficient. Because of these drawback, I am in the process of reimplementing many of the WEKA classifiers. As of now, only Naive Bayes has been reimplemented. In addition to Native Bayes, I have implemented a BPNB classifier[4].

4.3.1 BPNB

BPNB is a method developed by Dr. Leilei Chu. It is a Naive Bayes method that is based on relative probability of features.

4.4 Feature Selection

TODO

4.4.1 Text Filtering

Before the input text is split into features, it goes through heavy pre-processing. The text filtering involves normalizing the case, remove extra characters, and replacing special features.

Link Replacement

The first step in processing the text is to replace links. Following a link may provide information about a tweet, however the link text of the link itself provides no information. The presence of a link, however, can provide information about a tweet.

Twitter Specific Symbols

Tweets often contain several special symbols specific to tweets.

RT. “RT” stands for “re-tweet”. It means that the posted tweet is a repost of a tweet made by another user. This symbol contains no reference to the original post. “RT” usually appears at the beginning of the tweet. For example, after the comedian Conan O’Brien posted the following tweet:

If I’m ever a ghost, I hope the person I haunt has Netflix.

There were hundred of identical tweets that said:

RT: If I’m ever a ghost, I hope the person I haunt has Netflix.

#. In Twitter, a “#” (pronounced “hashtag”) is a reference to some topic in Twitter. Users can search for tweets by hashtag and see the collection of tweets supposedly about the same topic. A hashtag does not have to reference a pre-existing topic.

. An “@” in Twitter, simply pronounced “at”, is a reference to another Twitter user. A reference to a user will alert that user about the posted tweet. For example, the following tweet will reference my Twitter account.

Hi there, @eriquaugustine

Emoticon Parsing

Emoticons are parsed out and replaced with meta words. SPOONS emoticon parser was written by Ryan Hanarkis and Allen Dunlea. Emoticons provide a plethora of information about a tweet. Sarcasm aside, an emoticon can surmise the sentiment of an entire tweet.

Title Replacement

Because our tweets are always about Netflix, a television show and movie streaming service, titles are a common occurrence. However, movie and show title often contains words that can be detrimental to our analysis. For example, “Death At A Funeral” is the title of a movie, but contains two words that have very negative connotations: “death” and “funeral”.

Without title replacement, the following tweet would be very difficult to classify:

Death at a Funeral is hilarious! #netflix

However after title replacement, the tweet becomes very easy to classify:

;\$title\$ is hilarious! #netflix

Stemming

Stemming finds the root of a word. This allows words to be categorized by their roots which decreases the number of unique words being evaluated and

emphasizes linguistic patterns. This preprocessor uses Porter’s stemmer for the English language [16].

Stop Word Removal

Stopwords, or words that carry little or no semantic information, are identified based on a static table of words mapped to levels. Stopwords are assigned levels which allow processes to use different sets of stop words. All words less than 3 character are also automatically considered stop words.

Punctuation/Non-English Character Removal

Removes all punctuation and characters not in the English alphabet. This simplifies word extraction and comparison.

Meta Words

Below is an overview of meta words that SPOONS recognizes:

- $i\$link\$_i$ - Indicates the presence of a URL.
- $i\$emote:*\$_i$ - Replaces an emoticon.
- $i\$RT\$_i$ - Indicates that a tweet is a “retweet” (a repeat of another tweet).
- $i\$#\$_i$ - Inserted when a “hashtag” is found in a tweet.
- $i\$@\$_i$ - Inserted when a reference to another Twitter user is made.

4.5 Tweet Classes

Tweets generally fall into nine different categories:

- **Media** – Relate to a media story about Netflix.
- **Snafu** – Talk about a Netflix outage.
- **Complaint** – Complain about Netflix.
- **Happy** – Express joy about Netflix.
- **Neutral** – Neutral observation or comment about Netflix.
- **Watching** – Updates on what the user is currently watching.
- **Response** – Neutral response to another user in a Netflix-related conversation.
- **Refuse To Rate** – Used for tweets that we refuse to rate entirely (usually tweets that are in a different language than the training set).
- **Undetermined** – A default for all tweets that don't match any other class.

4.5.1 Tweet Groups

Because the goal of SPOONS is to detect anomalous times, it is useful to collapse the nine classes into three different groups that account for the different types of Netflix-related traffic.

- **Media**: Contains only the `media` class.
- **Bad**: Contains both the `outage` and `complaint` classes.




Figure 4.1: The different volumes for different tweet volumes during an outage (left) and media event (right).

- **Other/Normal:** Contains all other classes.

Figure 4.1 shows the amount of Netflix-related tweets during a Netflix outage and media event. During normal times, the normal traffic (purple line) is responsible for the majority of the overall traffic. However during outage and media events, we see that the bad (green line) and media (blue line) dominate the respective periods.

4.6 Training Set

The classifiers were trained on a small set of **759** tweets which were pulled from from periods of both normal and anomalous traffic. Each tweet in the training set was manually classified by multiple researchers until consensus about the classification was reached. Because the goal is anomalous traffic detection, the training set over-samples the tweets from **media**, **outage**, and **complaint** categories. Table 4.1 documents the structure of the training set and shows the number of tweets classified into each of the eight categories. Tweets were

Class	# Tweets	Class	# Tweets
Media	103	Neutral	66
Outage	158	Watching	135
Complaint	146	Response	30
Happy	147	Undetermined	48

Table 4.1: Overview of the Netflix-related Twitter post training set used to train classifiers in SPOONS.

allowed to belong to multiple classes because of posts like, “I love netflix! Watching Law and Order online!”, which could be classified as both happy and watching.

Chapter 5

Experimental Setup

5.1 Ground Truth

Netflix has provided us with a list of outages that occurred between March 14, 2011 and January 30, 2012. This list is not comprehensive and some of the times are questionable. Regardless, we use this as our base truth about all of the Netflix outages in that time period.

5.2 Success Metrics

The accuracy of outage detection is measured using three metrics:

- **Recall:** the percent of the reported events that were caught.
- **Precision:** the percent of the alerts generated that occurred during an outage event.

- **$F_{0.5}$ Measure:** a harmonic mean of precision and recall that weighs precision greater than recall.

The following definitions are used to calculate the metrics:

- **True Positive:** any intersection between a reported outage range and a detected outage range.
- **False Positive:** any detected outage that has no intersection in the events reported by Netflix.
- **False Negative:** no intersection on an event reported by Netflix is a false negative.

Netflix has specified that a precision of 0.5 is an acceptable amount of noise.

5.3 Classifier Evaluation

Each classifier is individually evaluated just on its ability to classify tweets against the training set. Each classifier varied three parameters: the type of filtering, the feature selection, and whether or not to collapse the classes into groups.

5.4 Outage Detection

The end goal of SPOONS is to be able to detect anomalous events. Therefore, each classifier is put into a full analysis pipeline and its ability to detect outages is evaluated. TODO(eriq): Pick monitor

5.4.1 Monitor Parameters

TODO(eriq): Numbers

Chapter 6

Results

TODO(eriq): Numbers

6.1 Classifier Evaluation

6.2 Outage Detection

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	48	50	0	1	2	2	0	0
neutral	0	1	40	3	8	8	2	0	4
snafu	0	0	48	69	3	6	32	0	0
watching	0	2	65	3	46	6	3	1	9
response	0	2	11	4	6	1	0	1	5
complaint	0	0	41	53	4	1	44	0	3
refuse to rate	0	2	24	0	2	2	0	18	0
happy	0	2	40	36	22	8	7	2	30

(a) Non-Weka NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	93	1	1	1	3	3	1	0
neutral	0	9	12	5	16	13	4	0	7
snafu	0	0	2	104	10	5	34	1	2
watching	0	4	23	5	83	7	2	1	10
response	0	3	7	4	9	0	0	1	6
complaint	0	6	3	66	6	3	57	0	5
refuse to rate	0	9	2	0	2	2	1	29	3
happy	0	4	17	41	32	6	10	2	35

(b) Non-Weka BPNBClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	85	4	1	2	4	0	6	1
neutral	0	0	25	3	11	13	3	0	11
snafu	0	0	6	94	7	7	25	1	18
watching	0	0	26	2	82	9	1	0	15
response	0	2	14	3	6	3	0	0	2
complaint	0	1	12	42	8	10	54	0	19
refuse to rate	0	6	4	0	1	2	2	30	3
happy	0	1	22	15	14	6	7	0	82

(c) NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	87	4	0	2	0	1	8	1
neutral	0	1	1	9	20	0	1	2	32
snafu	0	0	0	77	9	0	21	0	51
watching	0	0	0	6	105	1	1	0	22
response	0	2	0	4	9	0	1	0	14
complaint	0	2	0	36	10	0	24	0	74
refuse to rate	0	9	1	2	0	0	1	23	12
happy	0	1	0	26	21	0	2	0	97

(d) BayesNetClassifier

Table 6.1: Uncompressed, NGram, NoFilter Classification Results

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	95	1	0	2	0	1	3	1
neutral	0	4	9	5	13	1	7	3	24
snafu	0	0	5	84	3	0	36	2	28
watching	0	2	6	2	84	5	7	0	29
response	0	2	6	5	5	0	1	0	11
complaint	0	2	4	37	5	1	54	4	39
refuse to rate	0	11	2	4	0	1	6	14	10
happy	0	1	5	9	17	5	9	0	101

(e) J48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	56	0	0	1	2	0	0	44
neutral	0	0	3	2	7	4	5	0	45
snafu	0	0	0	45	0	4	59	0	50
watching	0	1	10	0	73	4	5	0	42
response	0	2	4	4	5	0	1	1	13
complaint	0	0	4	36	2	0	61	0	43
refuse to rate	0	2	0	0	1	1	0	4	40
happy	0	0	3	1	10	5	5	0	123

(f) KNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	98	0	0	2	2	0	0	1
neutral	0	3	11	2	16	7	7	0	20
snafu	0	0	1	87	4	4	41	0	21
watching	0	2	12	3	88	5	4	0	21
response	0	2	5	4	7	0	3	1	8
complaint	0	2	5	40	4	1	68	1	25
refuse to rate	0	9	1	1	1	1	0	18	17
happy	0	1	9	3	13	6	8	0	107

(g) SMOCClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	98	1	0	2	2	0	0	0
neutral	0	1	23	2	12	13	4	0	11
snafu	0	0	6	97	7	7	25	1	15
watching	0	0	21	2	92	5	2	0	13
response	0	2	14	3	6	3	0	0	2
complaint	0	1	11	45	10	8	53	0	18
refuse to rate	0	13	4	1	0	3	1	23	3
happy	0	1	17	17	19	6	9	0	78

(h) BinaryNaiveBayesClassifier

Table 6.1: Uncompressed, NGram, NoFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	97	0	2	0	0	1	2	1
neutral	0	2	4	10	10	0	3	1	36
snafu	0	0	0	61	6	0	26	0	65
watching	0	1	4	6	80	0	13	0	31
response	0	2	2	4	4	0	1	0	17
complaint	0	2	1	30	5	0	41	0	67
refuse to rate	0	14	0	2	1	0	5	5	21
happy	0	1	3	20	11	0	24	0	88

(i) BinaryJ48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	90	10	0	2	0	0	1	0
neutral	0	2	2	9	18	0	32	0	3
snafu	0	0	1	85	6	0	64	2	0
watching	0	0	1	4	107	0	20	0	3
response	0	2	1	4	9	0	14	0	0
complaint	0	1	2	33	10	0	99	0	1
refuse to rate	0	10	2	2	0	0	12	22	0
happy	0	1	2	26	19	0	91	0	8

(j) BinaryBayesNetClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	57	0	0	1	2	0	0	43
neutral	0	0	3	1	8	4	5	0	45
snafu	0	0	0	45	0	4	59	0	50
watching	0	1	10	0	74	4	5	0	41
response	0	2	5	4	5	0	0	1	13
complaint	0	0	4	36	2	0	61	0	43
refuse to rate	0	2	0	0	1	1	0	4	40
happy	0	0	3	1	10	5	5	0	123

(k) BinaryKNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	92	7	0	1	2	1	0	0
neutral	0	1	33	3	10	4	3	0	12
snafu	0	0	32	69	2	4	43	0	8
watching	0	2	28	2	84	4	2	0	13
response	0	2	11	4	5	0	0	1	7
complaint	0	1	38	37	3	0	55	0	12
refuse to rate	0	9	13	1	0	1	1	19	4
happy	0	1	50	1	13	5	2	0	75

(l) BinarySMOClassifier

Table 6.1: Uncompressed, NGram, NoFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	59	36	0	2	3	3	0	0
neutral	0	1	39	4	9	6	3	0	4
snafu	0	0	42	69	5	5	36	0	1
watching	0	2	64	5	45	5	4	0	10
response	0	2	11	4	5	1	1	1	5
complaint	0	1	39	51	4	1	47	0	3
refuse to rate	0	2	17	4	2	2	0	21	0
happy	0	2	39	13	21	5	7	1	59

(a) Non-Weka NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	88	2	4	2	3	3	1	0
neutral	0	9	9	5	19	7	8	0	9
snafu	0	1	2	100	9	5	39	0	2
watching	0	3	19	8	82	5	2	1	15
response	0	3	5	4	9	1	1	1	6
complaint	0	7	4	66	5	2	57	0	5
refuse to rate	0	7	1	4	5	2	1	26	2
happy	0	4	11	18	27	5	10	1	71

(b) Non-Weka BPNBClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	81	3	0	1	3	0	14	1
neutral	0	0	23	3	14	11	11	1	3
snafu	0	0	7	81	8	4	55	2	1
watching	0	0	26	0	82	5	12	0	10
response	0	2	12	3	6	3	1	0	3
complaint	0	1	9	36	9	8	77	3	3
refuse to rate	0	5	4	1	2	4	9	21	2
happy	0	0	18	5	22	5	31	1	65

(c) NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	94	1	0	0	0	1	5	2
neutral	0	3	0	8	3	1	49	0	2
snafu	0	0	0	48	1	2	103	0	4
watching	0	0	3	5	79	2	39	0	7
response	0	2	0	8	4	2	9	1	4
complaint	0	2	0	20	2	4	112	0	6
refuse to rate	0	14	0	5	0	0	20	8	1
happy	0	1	0	11	10	2	79	0	44

(d) BayesNetClassifier

Table 6.2: Uncompressed, NGram, EriqFilter Classification Results

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	99	0	0	0	0	2	1	1
neutral	0	4	10	2	8	3	13	1	25
snafu	0	0	5	71	4	3	50	0	25
watching	0	1	7	2	90	2	11	1	21
response	0	2	5	2	8	0	2	0	11
complaint	0	2	7	35	4	1	72	0	25
refuse to rate	0	13	2	1	3	1	6	10	12
happy	0	1	11	4	14	3	15	0	99

(e) J48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	58	17	9	1	2	7	0	9
neutral	0	0	21	13	8	4	10	0	10
snafu	0	0	20	57	0	4	72	0	5
watching	0	1	39	12	40	4	11	0	28
response	0	2	5	5	4	0	2	1	11
complaint	0	0	18	53	2	0	66	0	7
refuse to rate	0	0	13	9	2	1	7	4	12
happy	0	0	36	16	6	5	9	0	75

(f) KNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	98	1	0	1	2	0	1	0
neutral	0	4	7	3	15	5	6	0	26
snafu	0	0	3	86	2	4	50	0	13
watching	0	2	13	2	92	5	8	0	13
response	0	2	6	4	5	1	1	1	10
complaint	0	2	4	48	5	3	74	0	10
refuse to rate	0	11	0	1	2	2	6	8	18
happy	0	1	7	6	16	5	10	0	102

(g) SMOClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	95	2	0	1	2	0	2	1
neutral	0	2	23	4	14	9	11	0	3
snafu	0	0	7	85	8	4	52	1	1
watching	0	0	21	0	87	5	12	0	10
response	0	2	12	3	6	3	1	0	3
complaint	0	1	10	42	9	7	73	2	2
refuse to rate	0	9	4	1	2	4	9	17	2
happy	0	0	16	6	21	5	30	1	68

(h) BinaryNaiveBayesClassifier

Table 6.2: Uncompressed, NGram, EriqFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	99	0	2	0	0	0	0	2
neutral	0	4	0	25	13	0	22	0	2
snafu	0	0	0	97	6	0	55	0	0
watching	0	1	0	28	78	0	23	0	5
response	0	2	1	2	5	0	18	1	1
complaint	0	2	0	80	2	0	59	0	3
refuse to rate	0	14	0	14	3	0	12	4	1
happy	0	0	1	58	10	0	30	0	48

(i) BinaryJ48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	95	0	0	0	4	1	1	2
neutral	0	5	1	19	3	0	34	0	4
snafu	0	0	1	73	1	0	79	0	4
watching	0	0	1	10	75	0	32	0	17
response	0	2	1	6	4	0	11	0	6
complaint	0	2	1	27	2	0	107	0	7
refuse to rate	0	13	0	6	0	2	17	9	1
happy	0	0	2	18	9	0	63	0	55

(j) BinaryBayesNetClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	57	19	11	1	2	7	0	6
neutral	0	0	23	14	8	4	10	0	7
snafu	0	0	22	57	0	4	72	0	3
watching	0	1	41	12	42	4	11	0	24
response	0	2	7	7	4	0	1	1	8
complaint	0	0	18	57	3	0	66	0	2
refuse to rate	0	0	16	13	2	1	7	4	5
happy	0	0	40	17	6	5	9	0	70

(k) BinaryKNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	96	4	0	1	2	0	0	0
neutral	0	4	37	2	8	4	3	0	8
snafu	0	0	44	70	2	4	37	0	1
watching	0	2	32	1	84	5	4	0	7
response	0	2	11	4	6	0	0	1	6
complaint	0	2	43	38	4	0	57	0	2
refuse to rate	0	10	24	0	1	1	1	9	2
happy	0	1	46	1	13	5	3	0	78

(l) BinarySMOClassifier

Table 6.2: Uncompressed, NGram, EriqFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	82	10	0	5	4	0	1	1
neutral	0	2	17	5	18	18	2	1	3
snafu	0	6	10	99	13	7	20	1	2
watching	0	5	26	4	79	14	2	0	5
response	0	4	10	2	6	3	2	1	2
complaint	0	4	16	62	14	6	42	2	0
refuse to rate	0	18	6	2	1	5	0	16	0
happy	0	13	22	56	24	15	3	0	14

(a) Non-Weka NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	101	0	0	2	0	0	0	0
neutral	0	15	10	7	13	21	0	0	0
snafu	0	15	4	98	13	8	18	2	0
watching	0	10	13	6	92	13	1	0	0
response	0	11	7	4	4	4	0	0	0
complaint	0	17	8	59	11	13	38	0	0
refuse to rate	0	30	1	2	1	2	0	12	0
happy	0	18	12	56	36	7	2	0	16

(b) Non-Weka BPNBClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	66	3	1	2	6	0	24	1
neutral	1	0	20	2	11	20	3	0	9
snafu	0	0	7	96	4	10	24	0	17
watching	0	0	23	6	72	12	3	0	19
response	0	2	15	2	3	5	0	0	3
complaint	0	1	7	39	10	18	49	0	22
refuse to rate	0	4	6	4	0	3	1	29	1
happy	6	0	16	17	15	10	6	0	77

(c) NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	3	45	7	1	1	1	0	45	0
neutral	0	0	10	9	14	16	0	8	9
snafu	0	0	4	86	8	11	11	3	35
watching	0	0	10	8	81	16	0	0	20
response	0	2	6	4	5	6	0	5	2
complaint	0	0	11	42	7	14	16	4	52
refuse to rate	0	7	11	6	0	1	0	22	1
happy	0	0	7	37	20	7	2	7	67

(d) BayesNetClassifier

Table 6.3: Uncompressed, TweetFSG, NoFilter Classification Results

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	97	0	0	1	0	1	2	2
neutral	0	3	12	7	14	4	9	0	17
snafu	0	0	5	89	2	1	41	2	18
watching	0	1	10	7	90	1	6	1	19
response	0	2	6	6	5	0	4	1	6
complaint	0	2	3	42	9	2	61	3	24
refuse to rate	0	12	1	6	1	1	8	16	3
happy	0	1	8	11	22	2	20	1	82

(e) J48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	81	0	0	3	2	1	13	3
neutral	0	7	4	4	11	4	7	15	14
snafu	0	2	2	45	4	4	53	19	29
watching	0	6	10	3	80	4	7	9	16
response	0	3	4	5	5	0	2	4	7
complaint	0	6	3	37	9	1	50	15	25
refuse to rate	0	7	0	0	1	2	0	33	5
happy	0	3	5	5	15	7	10	13	89

(f) KNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	94	0	0	2	2	1	4	0
neutral	0	1	14	2	11	6	9	2	21
snafu	0	0	3	92	3	4	40	0	16
watching	0	2	16	1	84	7	4	0	21
response	0	2	9	4	5	0	0	1	9
complaint	0	1	7	44	7	1	64	2	20
refuse to rate	0	9	2	0	1	2	1	28	5
happy	0	1	10	2	17	6	7	1	103

(g) SMOClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	84	3	2	2	1	0	10	1
neutral	1	0	20	2	14	18	3	0	8
snafu	0	0	7	99	5	8	23	0	16
watching	0	0	19	7	85	9	2	0	13
response	0	2	15	2	4	4	0	0	3
complaint	0	1	7	41	11	17	46	0	23
refuse to rate	0	10	5	5	0	3	1	23	1
happy	6	0	14	19	22	9	6	0	71

(h) BinaryNaiveBayesClassifier

Table 6.3: Uncompressed, TweetFSG, NoFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	95	2	1	1	0	1	0	3
neutral	0	2	5	10	18	0	13	1	17
snafu	0	0	3	99	4	0	36	4	12
watching	0	1	7	14	96	0	3	0	14
response	0	2	2	6	9	0	9	0	2
complaint	0	2	2	50	6	0	52	4	30
refuse to rate	0	14	0	11	0	0	8	11	4
happy	0	1	5	33	25	0	13	1	69

(i) BinaryJ48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	78	8	0	1	1	0	15	0
neutral	0	2	6	12	19	2	0	18	7
snafu	0	0	1	91	2	2	7	15	40
watching	0	0	2	10	101	1	1	6	14
response	0	3	3	6	8	1	0	8	1
complaint	0	2	6	41	7	3	20	19	48
refuse to rate	0	9	3	5	0	3	0	27	1
happy	0	0	7	38	24	0	0	11	67

(j) BinaryBayesNetClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	81	0	2	3	2	1	11	3
neutral	0	8	4	5	14	5	6	14	10
snafu	0	3	3	52	13	5	51	17	14
watching	0	7	12	2	87	5	6	6	10
response	0	3	5	5	5	0	1	4	7
complaint	0	7	4	40	10	4	49	11	21
refuse to rate	0	7	0	1	3	3	0	31	3
happy	0	3	6	9	23	7	9	9	81

(k) BinaryKNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	89	9	0	1	2	1	1	0
neutral	0	1	31	1	8	5	8	0	12
snafu	0	0	38	64	1	4	44	0	7
watching	0	2	33	1	80	5	4	0	10
response	0	2	10	4	5	0	1	1	7
complaint	0	1	35	38	4	0	57	0	11
refuse to rate	0	10	8	0	1	1	1	25	2
happy	0	1	45	0	17	4	1	0	79

(l) BinarySMOClassifier

Table 6.3: Uncompressed, TweetFSG, NoFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	93	5	0	0	4	0	1	0
neutral	0	8	16	6	14	18	1	3	0
snafu	0	8	6	100	17	7	15	5	0
watching	0	10	16	5	74	23	3	1	3
response	0	6	8	1	3	8	0	3	1
complaint	0	12	9	57	9	12	43	3	1
refuse to rate	0	30	4	0	1	9	0	3	1
happy	0	17	9	32	28	19	1	1	40

(a) Non-Weka NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	102	0	0	1	0	0	0	0
neutral	0	25	8	7	9	15	2	0	0
snafu	0	21	1	99	10	7	20	0	0
watching	0	29	9	9	78	9	0	0	1
response	0	16	2	1	4	7	0	0	0
complaint	0	29	2	59	6	10	40	0	0
refuse to rate	0	40	1	1	2	3	0	0	1
happy	0	23	6	38	26	9	1	0	44

(b) Non-Weka BPNBClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	66	5	3	1	3	0	24	1
neutral	1	1	24	9	9	10	4	3	5
snafu	0	0	6	91	7	10	27	1	16
watching	0	0	21	11	63	8	2	2	28
response	0	2	14	3	2	5	0	2	2
complaint	0	1	10	44	10	12	50	7	12
refuse to rate	0	4	5	5	5	5	0	23	1
happy	6	0	15	19	19	6	5	1	76

(c) NaiveBayesClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	2	29	1	0	3	1	3	64	0
neutral	0	0	1	10	14	9	2	21	9
snafu	0	0	3	65	10	13	22	7	38
watching	0	0	7	13	55	14	3	10	33
response	0	0	1	3	4	11	0	9	2
complaint	0	0	3	32	11	13	20	15	52
refuse to rate	0	5	4	5	5	8	0	20	1
happy	0	0	1	19	25	8	11	10	73

(d) BayesNetClassifier

Table 6.4: Uncompressed, TweetFSG, EriqFilter Classification Results

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	99	0	1	0	0	1	1	1
neutral	0	4	11	7	11	5	8	4	16
snafu	0	0	1	77	6	1	48	3	22
watching	0	1	10	7	87	0	9	2	19
response	0	2	6	3	6	1	2	0	10
complaint	0	2	8	42	7	2	59	2	24
refuse to rate	0	13	4	7	1	0	7	13	3
happy	0	1	7	16	15	3	8	4	93

(e) J48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	83	0	1	3	2	0	14	0
neutral	0	8	3	2	9	4	3	24	13
snafu	0	1	0	43	10	3	44	36	21
watching	0	2	9	7	59	5	5	29	19
response	0	3	4	3	5	0	2	6	7
complaint	0	3	2	35	5	0	45	29	27
refuse to rate	0	13	0	2	1	1	1	27	3
happy	0	3	4	10	15	5	3	30	77

(f) KNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	95	1	0	1	2	1	3	0
neutral	0	3	13	3	9	6	7	4	21
snafu	0	0	4	80	2	4	47	6	15
watching	0	2	13	7	85	6	3	4	15
response	0	2	5	5	5	0	3	4	6
complaint	0	2	4	43	7	2	65	9	14
refuse to rate	0	11	2	3	1	2	1	23	5
happy	0	1	10	3	14	7	9	2	101

(g) SMOClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	78	3	3	2	3	0	13	1
neutral	1	2	23	9	10	10	4	2	5
snafu	0	0	6	97	7	9	22	1	16
watching	0	0	18	15	75	4	2	2	19
response	0	2	14	4	2	5	0	2	1
complaint	0	1	10	44	10	12	50	7	12
refuse to rate	0	9	4	5	5	5	0	19	1
happy	1	0	15	20	20	5	3	1	82

(h) BinaryNaiveBayesClassifier

Table 6.4: Uncompressed, TweetFSG, EriqFilter Classification Results Cont.

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	99	0	0	0	0	1	0	3
neutral	0	4	0	15	11	0	19	0	17
snafu	0	0	0	102	7	0	32	0	17
watching	0	1	0	19	82	0	16	0	17
response	0	2	1	3	6	0	9	0	9
complaint	0	2	1	63	5	0	54	0	21
refuse to rate	0	13	0	4	1	0	12	9	9
happy	0	0	1	44	11	0	27	0	64

(i) BinaryJ48Classifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	74	3	0	0	0	0	26	0
neutral	0	5	0	10	4	2	1	33	11
snafu	0	0	0	70	1	2	8	22	55
watching	0	0	1	13	66	2	0	22	31
response	0	3	0	4	4	4	0	13	2
complaint	0	3	0	22	2	3	20	33	63
refuse to rate	0	11	0	4	0	2	0	29	2
happy	0	0	0	19	8	2	3	26	89

(j) BinaryBayesNetClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	83	0	1	4	2	0	13	0
neutral	0	8	3	3	13	4	3	21	11
snafu	0	1	0	52	16	4	44	24	17
watching	0	4	9	6	68	5	4	24	15
response	0	3	5	3	5	0	1	6	7
complaint	0	5	2	36	12	0	44	25	22
refuse to rate	0	14	0	2	1	1	0	27	3
happy	0	3	4	14	22	5	3	29	67

(k) BinaryKNNClassifier

	undecided	media	neutral	snafu	watching	response	complaint	refuse to rate	happy
undecided	0	0	0	0	0	0	0	0	0
media	0	89	9	0	1	2	1	1	0
neutral	0	3	39	1	8	4	5	0	6
snafu	0	0	39	68	1	3	45	0	2
watching	0	2	33	2	84	3	3	0	8
response	0	2	12	3	5	0	1	2	5
complaint	0	2	47	31	3	0	58	4	1
refuse to rate	0	10	17	0	1	1	1	16	2
happy	0	1	62	1	9	5	4	0	65

(l) BinarySMOClassifier

Table 6.4: Uncompressed, TweetFSG, EriqFilter Classification Results Cont.

	media	snafu	other
media	48	2	53
snafu	0	198	106
other	9	58	359

(a) Non-Weka NaiveBayesClassifier

	media	snafu	other
media	93	4	6
snafu	6	261	37
other	29	72	325

(b) Non-Weka BPNBClassifier

	media	snafu	other
media	85	1	17
snafu	1	215	88
other	9	36	381

(c) NaiveBayesClassifier

	media	snafu	other
media	87	1	15
snafu	2	158	144
other	13	53	360

(d) BayesNetClassifier

	media	snafu	other
media	95	1	7
snafu	2	211	91
other	20	55	351

(e) J48Classifier

	media	snafu	other
media	56	0	47
snafu	0	201	103
other	5	23	398

(f) KNNClassifier

	media	snafu	other
media	98	0	5
snafu	2	236	66
other	17	35	374

(g) SMOClassifier

	media	snafu	other
media	98	0	5
snafu	1	220	83
other	17	41	368

(h) BinaryNaiveBayesClassifier

	media	snafu	other
media	97	3	3
snafu	2	158	144
other	20	88	318

(i) BinaryJ48Classifier

	media	snafu	other
media	90	0	13
snafu	1	281	22
other	15	214	197

(j) BinaryBayesNetClassifier

	media	snafu	other
media	57	0	46
snafu	0	201	103
other	5	21	400

(k) BinaryKNNClassifier

	media	snafu	other
media	92	1	10
snafu	1	204	99
other	15	19	392

(l) BinarySMOClassifier

Table 6.5: Compressed, NGram, NoFilter Classification Results

	media	snafu	other
media	59	3	41
snafu	1	203	100
other	9	45	372

(a) Non-Weka NaiveBayesClassifier

	media	snafu	other
media	88	7	8
snafu	8	262	34
other	26	61	339

(b) Non-Weka BPNBClassifier

	media	snafu	other
media	81	0	22
snafu	1	249	54
other	7	76	343

(c) NaiveBayesClassifier

	media	snafu	other
media	94	1	8
snafu	2	283	19
other	20	233	173

(d) BayesNetClassifier

	media	snafu	other
media	99	2	2
snafu	2	228	74
other	21	58	347

(e) J48Classifier

	media	snafu	other
media	58	16	29
snafu	0	248	56
other	3	94	329

(f) KNNClassifier

	media	snafu	other
media	98	0	5
snafu	2	258	44
other	20	47	359

(g) SMOClassifier

	media	snafu	other
media	95	0	8
snafu	1	252	51
other	13	77	336

(h) BinaryNaiveBayesClassifier

	media	snafu	other
media	99	2	2
snafu	2	291	11
other	21	232	173

(i) BinaryJ48Classifier

	media	snafu	other
media	95	1	7
snafu	2	286	16
other	20	216	190

(j) BinaryBayesNetClassifier

	media	snafu	other
media	57	18	28
snafu	0	252	52
other	3	101	322

(k) BinaryKNNClassifier

	media	snafu	other
media	96	0	7
snafu	2	202	100
other	19	19	388

(l) BinarySMOClassifier

Table 6.6: Compressed, NGram, EriqFilter Classification Results

	media	snafu	other
media	82	0	21
snafu	10	223	71
other	42	78	306

(a) Non-Weka NaiveBayesClassifier

	media	snafu	other
media	101	0	2
snafu	32	213	59
other	84	78	264

(b) Non-Weka BPNBClassifier

	media	snafu	other
media	66	1	36
snafu	1	208	95
other	6	44	376

(c) NaiveBayesClassifier

	media	snafu	other
media	45	1	57
snafu	0	155	149
other	9	66	351

(d) BayesNetClassifier

	media	snafu	other
media	97	1	5
snafu	2	233	69
other	19	84	323

(e) J48Classifier

	media	snafu	other
media	81	1	21
snafu	8	185	111
other	26	43	357

(f) KNNClassifier

	media	snafu	other
media	94	1	8
snafu	1	240	63
other	15	30	381

(g) SMOClassifier

	media	snafu	other
media	84	2	17
snafu	1	209	94
other	12	47	367

(h) BinaryNaiveBayesClassifier

	media	snafu	other
media	95	2	6
snafu	2	237	65
other	20	120	286

(i) BinaryJ48Classifier

	media	snafu	other
media	78	0	25
snafu	2	159	143
other	14	72	340

(j) BinaryBayesNetClassifier

	media	snafu	other
media	81	3	19
snafu	10	192	102
other	28	44	354

(k) BinaryKNNClassifier

	media	snafu	other
media	89	1	13
snafu	1	203	100
other	16	21	389

(l) BinarySMOClassifier

Table 6.7: Compressed, TweetFSG, NoFilter Classification Results

	media	snafu	other
media	93	0	10
snafu	20	215	69
other	71	49	306

(a) Non-Weka NaiveBayesClassifier

	media	snafu	other
media	102	0	1
snafu	50	218	36
other	133	59	234

(b) Non-Weka BPNBClassifier

	media	snafu	other
media	66	3	34
snafu	1	212	91
other	7	58	361

(c) NaiveBayesClassifier

	media	snafu	other
media	29	3	71
snafu	0	139	165
other	5	66	355

(d) BayesNetClassifier

	media	snafu	other
media	99	2	2
snafu	2	226	76
other	21	74	331

(e) J48Classifier

	media	snafu	other
media	83	1	19
snafu	4	167	133
other	29	38	359

(f) KNNClassifier

	media	snafu	other
media	95	1	7
snafu	2	235	67
other	19	44	363

(g) SMOClassifier

	media	snafu	other
media	78	3	22
snafu	1	213	90
other	13	62	351

(h) BinaryNaiveBayesClassifier

	media	snafu	other
media	99	1	3
snafu	2	251	51
other	20	168	238

(i) BinaryJ48Classifier

	media	snafu	other
media	74	0	29
snafu	3	120	181
other	19	54	353

(j) BinaryBayesNetClassifier

	media	snafu	other
media	83	1	19
snafu	6	176	122
other	32	39	355

(k) BinaryKNNClassifier

	media	snafu	other
media	89	1	13
snafu	2	202	100
other	18	21	387

(l) BinarySMOClassifier

Table 6.8: Compressed, TweetFSG, EriqFilter Classification Results

Chapter 7

Conclusions

7.1 Current Limitations of SPOONS

Severity Nature of Outage Malicious Tweet Attack Know What To Search
For (dynamic search generation)

7.2 Current and Future Work

Brett - Feasability of SPOONS as a comercial multi-target, source system.
All sorts of stuff Open Source System Dynamic Training Set

7.2.1 Advanced Sentiment Analysis

Kim - Advanced Sentiment Analysis

Glossary

AWS Amazon Web Services. Cloud computing offered by Amazon.

EC2 Elastic Compute Cloud. Instance based cloud computing machines offered through AWS.

Netflix Inc. [NASDAQ: NFLX] is the world's leading Internet subscription service for enjoying movies and TV series with more than 23 million streaming members in the United States, Canada, Latin America, the United Kingdom and Ireland[?].

Real Time Some of Netflix's services stream to customers in real time which means the users expect to get immediate responses from those services. So when they go down, the customers want the problem to be fixed immediately. These analysis methods need to have real time responses that are as close to immediate detection as possible. This means that the system needs to use whatever information it has available to it up to right before the outage to detect the event and alert Netflix engineers.

SPOONS Swift Perception Of Online Negative Situations. The name of the system presented in this paper.

Time Series Analysis The analysis of a series of data points over time. In this work those data points are the volume or estimated sentiment of a subset of the traffic about Netflix on Twitter during a time period.

Tweet A micro-post to a Twitter service. Tweets are limited to 140 characters.

Twitter Twitter is a social media service that allows users to post tweets (micro-posts) about any topic.

Bibliography

- [1] Innodb table and index structures.
- [2] E. Augustine, C. Cushing, A. Dekhtyar, M. Tognetti, and K. Paterson. Outage detection via real-time social stream analysis: Leveraging the power of online complaints. In *WWW 2012: Proceedings of the 21st World Wide Web Conference*. ACM, 2012.
- [3] N. Bansal and N. Koudas. Blogscope: a system for online analysis of high volume text streams. In *Proceedings of the 33rd international conference on Very large data bases, VLDB '07*, pages 1410–1413. VLDB Endowment, 2007.
- [4] L. Chu. Research on chinese text categorization method oriented to imbalanced corpus. may 2012.
- [5] A. Culotta. Detecting influenza outbreaks by analyzing twitter messages. In *KDD Workshop on Social Media Analytics*, 2010.
- [6] C. Cushing. Detecting netflix service outages through analysis of twitter posts. Master’s thesis, California Polytechnic State University - San Luis Obispo, june 2012.
- [7] M. Grinev, M. Grineva, A. Boldakov, L. Novak, A. Syssoev, and D. Lizorkin.

- Sifting micro-blogging stream for events of user interest. In *Proceedings of the 32nd international ACM SIGIR conference on Research and development in information retrieval*, SIGIR '09, pages 837–837, New York, NY, USA, 2009. ACM.
- [8] M. Hall, E. Frank, G. Holmes, B. Pfahringer, P. Reutemann, and I. H. Witten. The weka data mining software: an update. *SIGKDD Explor. Newsl.*, 11(1):10–18, 2009.
- [9] L. Hong and B. D. Davison. Empirical study of topic modeling in twitter. In *Proceedings of the First Workshop on Social Media Analytics*, SOMA '10, pages 80–88, New York, NY, USA, 2010. ACM.
- [10] J. Huang, B. Zhou, Q. Wu, X. Wang, and Y. Jia. Contextual correlation based thread detection in short text message streams. *J. Intell. Inf. Syst.*, 38(2):449–464, Apr. 2012.
- [11] F. Jabr. Using twitter to follow trends beats the stock market. *NewScientist*, (2829), Sept. 2011. <http://www.newscientist.com/article/mg21128295.900-using-twitter-to-follow-trends-beats-the-stock-market.html>.
- [12] Y. Matsu, M. Okazaki, and T. Sakak. Earthquake shakes twitter users: real-time event detection by social sensors. In *WWW 2010: Proceedings of the 19th World Wide Web Conference*, 2010. <http://ymatsuo.com/papers/www2010.pdf>.
- [13] K. McEntee. personal communication, 2011.
- [14] B. Pang, L. Lee, and S. Vaithyanathan. Thumbs up?: sentiment classification using machine learning techniques. In *Proceedings of the ACL-02 conference on Empirical methods in natural language processing - Volume*

- 10, EMNLP '02, pages 79–86, Stroudsburg, PA, USA, 2002. Association for Computational Linguistics.
- [15] J. C. Platt. Advances in kernel methods. chapter Fast training of support vector machines using sequential minimal optimization, pages 185–208. MIT Press, Cambridge, MA, USA, 1999.
 - [16] M. F. Porter. An algorithm for suffix stripping. In K. Sparck Jones and P. Willett, editors, *Readings in information retrieval*, pages 313–316. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1997. <http://tartarus.org/martin/PorterStemmer>.
 - [17] J. R. Quinlan. *C4.5: programs for machine learning*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1993.
 - [18] Twitter. #numbers, Mar. 2011. <http://blog.twitter.com/2011/03/numbers.html>.
 - [19] Twitter. Terms of service, June 2011. <https://twitter.com/tos>.
 - [20] Twitter. Get search/tweets, Oct. 2012. <https://dev.twitter.com/docs/api/1.1/get/search/tweets>.